

# Assignment 3 - CS 6240

- Sriharsha Srinivasa Karthik Kaipa, Sec 01

## Running the code:

The source code for this assignment is available in the folder named "assignment3". The "src" folder contains the final code. Make sure the inputs are in the folder "input" alongside the "src" folder. Run instructions for the code are as follows:

- For cleaning the project, building the jar and then running it from the command line, please use the following rules:  
~\$ make pagerank
- For cleaning the project, use the following rule  
make clean
- For building the jar, use the following rule  
make jar

## Pseudo Code:

### Pre-processing:

Used code provided by professor. No changes.

### PageRank Mapper:

#### **Mapper Task:**

```
map (line) {  
    rank = rank + 0.85 * DELTA / COUNT;  
    emit(page, "!"); // realpage  
    if adjacencyList is empty {  
        DELTA += (rank * 100.0);  
    }  
    for all lint in adjacencyList {  
        emit(link, {rank, adjacencyList.length});  
    }  
    COUNT += 1;  
    emit(page, {!, adjacencyList});  
}
```

PageRank Reducer:

**Reducer Task:**

```
reduce (page, list = ["!", {r1, n1}, ... , {l, adL}]) {
    for item in list {
        if item is "!" {
            realPage = true;
            continue;
        }
        if item is {l, adL} {
            adjList = adL;
            continue;
        }
        sumPageRank += ri / ni;
    }
    if not realPage
        return;
    rank = 0.85 * sumPageRank + 0.15 / COUNT;
    emit(page, {rank, adjList});
}
```

TopK Mapper:

**Mapper Task:**

```
setup () {
    HashMap hs = new HashMap();
}
map (line) {
    rank = rank + 0.85 * DELTA / COUNT;
    hs.put(page, rank);
}
cleanup () {
    hs = sortByValue(hs);
    count = 0;
    for each {page, rank} in hs {
        emit(null, {page, rank});
        count++;
        if count >= k
            break;
    }
}
```

TopK Reducer:

**Reducer Task:**

```
reduce (null, list = [{p1, r1}, {p2, r2}...]) {
    HashMap hs = new HashMap();
    for each {page, rank} in list {
        hs.put(page, rank);
    }
    hs = sortByValue(hs);
    count = 0;
    for each {page, rank} in hs {
        emit(null, {page, rank});
        count++;
        if count >= k
            break;
    }
}
```

Data transfer for 1 master 10 core run (all iterations):

**Iteration 1:**

Mappers -> Reducers:

Map output materialized bytes=1410334471 (~1.41 GB)

Reducers -> HDFS:

S3: Number of bytes written=1398201399 (~1.398 GB)

**Iteration 2:**

Mappers -> Reducers:

Map output materialized bytes=2245224243 (~2.245 GB)

Reducers -> HDFS:

S3: Number of bytes written=1398676946 (~1.399 GB)

**Iteration 3:**

Mappers -> Reducers:

Map output materialized bytes=2253622193 (~2.254 GB)

Reducers -> HDFS:

S3: Number of bytes written=1398604011 (~1.399 GB)

**Iteration 4:**

Mappers -> Reducers:

Map output materialized bytes=2257397797 (~2.257 GB)

Reducers -> HDFS:

S3: Number of bytes written=1398553684 (~1.398 GB)

**Iteration 5:**

Mappers -> Reducers:

Map output materialized bytes=2262446742 (~2.262 GB)

Reducers -> HDFS:

S3: Number of bytes written=1398519492 (~1.398 GB)

#### **Iteration 6:**

Mappers -> Reducers:

Map output materialized bytes=2265690722 (~2.266 GB)

Reducers -> HDFS:

S3: Number of bytes written=1398486868 (~1.398 GB)

#### **Iteration 7:**

Mappers -> Reducers:

Map output materialized bytes=2265424162 (~2.265 GB)

Reducers -> HDFS:

S3: Number of bytes written=1398445510 (~1.398 GB)

#### **Iteration 8:**

Mappers -> Reducers:

Map output materialized bytes=2260843034 (~2.261 GB)

Reducers -> HDFS:

S3: Number of bytes written=1398426886 (~1.398 GB)

#### **Iteration 9:**

Mappers -> Reducers:

Map output materialized bytes=2251275329 (~2.251 GB)

Reducers -> HDFS:

S3: Number of bytes written=1398411854 (~1.398 GB)

#### **Iteration 10:**

Mappers -> Reducers:

Map output materialized bytes=2250241015 (~2.25 GB)

Reducers -> HDFS:

S3: Number of bytes written=1398375929 (~1.398 GB)

The amount of data transferred from Mapper to Reducer in the first run is significantly lower in comparison to the rest of the runs. This could be because in the very first run, all the page ranks for all the pages is set to the string "NaN" while in later iterations, it varies from repeated calculation.

The amount of data transferred from Reducer to file remains more or less constant over repeated runs. I don't have an explanation for why this is happening.

### Performance Comparison:

#### **6 m4.large machines:**

Preprocessing time: 33:26 (mm:ss)

PageRank run time: 34:18 (mm:ss)

TopK run time: 01:05 (mm:ss)

#### **10 m4.large machines:**

Preprocessing time: 23:46 (mm:ss)

PageRank run time: 27:32 (mm:ss)

TopK run time: 00:45 (mm:ss)

The expected behavior was that every phase except the TopK run would be almost twice as fast with 11 machines when compared to 6 machines. As can be seen from the run times above, it is the case. The only phase which doesn't show twice the speedup is the TopK run. This is happening because unlike the preprocessing and pagerank phases, this phase only has only one reduce task running, irrespective of the number of core machines dedicated to MapReduce. There is a decent amount of speedup, however, but the speedup is not as significant as that seen in the other phases.

**Sample Wiki PageRank:**

0.010656700141346977	United_States_09d4
0.008220639776057527	Wikimedia_Commons_7b57
0.0064952402353202566	Country
0.004711283065929041	England
0.0045635853871618965	Germany
0.004414657920218037	United_Kingdom_5ad7
0.004380909397817181	Europe
0.004359713264360057	France
0.004308947196780668	Water
0.004172974263868258	Animal
0.004124838763467	City
0.003990314502819405	Earth
0.0032022623206404612	Wiktionary
0.003190687587569432	Asia
0.0031397724967365728	Week
0.0030900864024562403	Money
0.0030672136089680358	Plant
0.0030177594605749116	Computer
0.003006232620767282	Sunday
0.002965622101995127	Monday
0.0029584070412357914	English_language
0.002938916857904719	Italy
0.0029376069177789255	Wednesday
0.002891313724770175	India
0.0028679582384330796	Friday
0.002859088683119939	Government
0.0028359297174876197	Saturday
0.002800200551367329	Thursday
0.0027797327121635543	Tuesday
0.0027158710035411282	Spain
0.002628420081539975	Number
0.002623675284060614	Japan
0.0025578671694649027	Canada
0.002441746192137675	People
0.002405821692748891	Human
0.0023446916812946664	Australia

0.002326782767279969	China
0.002308180786288308	Day
0.0023064364423891173	Wikimedia_Foundation_83d9
0.002241928372616844	Food
0.002174809526271151	Energy
0.0021670058929277227	Mathematics
0.0021543794871223734	Science
0.002136262375946121	Television
0.0021041390580040977	Sun
0.0020360651170750724	Capital_(city)
0.0020252283756537004	Music
0.0020069998433729277	Russia
0.0019582379703325797	State
0.0019081671466966736	Greece
0.0019024539949288237	index
0.0019005915011195112	Scotland
0.0018930928409946729	Year
0.001891613493292557	Language
0.0018364339229566622	Metal
0.0018346998701719212	2004
0.0017946652579968441	Wikipedia
0.0017803392350589503	London
0.0017796319188465941	Greek_language
0.0017453826941465605	Religion
0.0017393334979107248	Sound
0.0016949247662825155	Africa
0.0016665522626982648	Planet
0.0016394780227830278	20th_century
0.001634741537357645	Poland
0.0016095159378284234	19th_century
0.0016014789732554573	Law
0.0015636273637329805	Geography
0.0015584184765731562	Liquid
0.0015508915158835817	World
0.001530950799436273	Scientist
0.00150874870954171	Society
0.001492023174735682	Inhabitant
0.0014871899132532646	Latin
0.0014829147468236852	Netherlands
0.0014804876270619798	Sweden
0.0014783865575128813	War
0.0014709528105714936	History
0.0014363723901016283	Light
0.001436325981133582	Atom
0.0014225398458921927	Building
0.001417104455261456	Culture

0.0014159428457179684	God
0.0013843277505049129	Centuries
0.0013723761053671705	Information
0.001371173133459838	Capital_city
0.0013708669728049654	Turkey
0.0013551765855783241	Plural
0.00134247735034764	Portugal
0.0013353346242485215	Chemical_element
0.0013169168937867152	University
0.0013154875173614937	Denmark
0.0013129632257306593	Book
0.001312082444868015	Austria
0.0013042374491607416	Species
0.0012939934363504517	Disease
0.0012786131678996253	North_America_e7c4
0.0012645015761637707	Biology
0.0012585104653527607	Ocean
0.0012529807724905777	U.S._state_5a68

#### **Full Wiki PageRank:**

0.020328173945099632	United_States_09d4
0.01676857403516821	2006
0.009379956797138371	United_Kingdom_5ad7
0.007976707136631527	2005
0.007816284494847574	Biography
0.006403586799651899	Canada
0.006328172721341767	England
0.005875132531300701	France
0.005623137574084618	2004
0.005465648295730867	Geographic_coordinate_system
0.005193700055033483	Australia
0.005163551443629481	Germany
0.004632736551948406	India
0.004548057304136969	2003
0.00442459242921374	Japan
0.00387304648328624	Internet_Movie_Database_7ea7
0.0037118203429510376	2001
0.0036294043486943027	2002
0.0035209352107723737	Italy
0.0034763517281045696	Record_label
0.003460132535920062	2000
0.0034325887639080664	Europe
0.0032467589437966725	Population_density
0.003197766312305399	World_War_II_d045
0.0031443132743244552	London
0.003031902837474414	1999

0.0029645902411086045	Music_genre
0.00293111443324066	Spain
0.0027924565089871667	Race_(United_States_Census)_a07d
0.002767867474258461	English_language
0.0027088960240600193	Football_(soccer)
0.0026971219995995544	Russia
0.0026310676898812437	Wiktionary
0.002629982966837057	1998
0.002490265453607359	1997
0.002480644449033358	Scotland
0.00244392496105335	Television
0.002399543540459641	Census
0.0023872836566535315	New_York_City_1428
0.0023658860548404536	Sweden
0.0023461114805196906	Wikimedia_Commons_7b57
0.0023417826815513455	1996
0.00229899361301265	Square_mile
0.0022564758736645365	California
0.0021948342750785434	1995
0.002158589761611589	New_Zealand_2311
0.002133590863459327	Scientific_classification
0.002114859056100764	China
0.0020939112947327577	Actor
0.0020917731177998266	1994
0.002049047181290158	Netherlands
0.0020474821402696455	Public_domain
0.002041315387841829	Politician
0.0020032764602457664	Film
0.0020006233317782663	New_York_3da4
0.0019885924252686584	United_States_Census_Bureau_2c85
0.0019836226086137375	1993
0.0019644005879764937	1991
0.001947383042797961	1990
0.0019397407028829239	Norway
0.0019302329005678675	Album
0.0018940257913438161	1992
0.0018881123586562952	Marriage
0.0018878081770618937	Population
0.0018662397823301633	Ireland
0.0018657206747686049	Poland
0.0018608058759947799	Record_producer
0.0018358869318855484	Personal_name
0.0017978074335498295	Per_capita_income
0.0017889129420747539	Brazil
0.0017607738686730645	1989
0.0017199458511551401	Mexico



0.0017064644912262938	1980
0.0016905682817106625	Studio_album
0.001663916076440879	1986
0.0016612547348724238	Poverty_line
0.0016330279890918664	1985
0.001626331837988042	1982
0.0016174863282596252	1981
0.0016120809203950246	Animal
0.0016092398801357975	1979
0.0016079671463468603	South_Africa_1287
0.001603195334692069	1987
0.0016001705386681042	1983
0.0015993641704791716	1984
0.0015740804855251243	1974
0.00157222768911212	Switzerland
0.0015681051082448592	1988
0.001553505183337771	January_1
0.0015281069107824495	Latin
0.0015218300236394188	1970
0.0015158575159968366	1976
0.0014971983401575616	1975
0.0014851192435254832	Paris
0.0014725613132124326	French_language
0.0014718499399261108	1978
0.001470968169533231	Greece
0.0014685829069004885	1977
0.0014648182391194186	1972
0.0014613959487998417	1969

The ranking seems reasonable for the full Wiki page dump. The same is more or less applicable for sample data as well, except for the page named "index", which shouldn't be there.